

## WEEK 1

machine learning definition:

Field of study that gives computers the ability to learn without being explicitly programmed.

—Arthur Samuel

A computer is said to learn from experience  $E$  with respect to some class of tasks  $T$  and performance  $P$ , if its performance at tasks in  $T$ , as measured by  $P$ , improves with experience  $E$ .

—Tom Mitchell


Intro to supervised learning

"right answers" given  
categorized as:

- { Regression: Predict continuous valued output
- { Classification: Discrete valued output

Intro to unsupervised learning

Approach problems with little or no ideas  
about what the results should look like




Supervised learning model


Training set

linear regression with one variable:

$$h_0(x) = \theta_0 + \theta_1 x$$



Cost function



$$\theta_0, \theta_1 \text{ to minimize: } \frac{1}{2n} \sum_{i=1}^n (h_0(x^{(i)}) - y^{(i)})^2$$

$n$ : # training examples

$$h_0(x^{(i)}) = \theta_0 + \theta_1 x^{(i)}$$

$y$ : actual value

$$\text{Cost function: } J(\theta_0, \theta_1) = \frac{1}{2n} \sum_{i=1}^n (h_0(x^{(i)}) - y^{(i)})^2$$

(Squared error function)

Commonly used for regression problems

Intuition:

$$\text{Hypothesis: } h_0(x) = \theta_0 + \theta_1 x$$

Parameters:  $\theta_0, \theta_1$

$$\text{Cost Function: } J(\theta_0, \theta_1) = \frac{1}{2n} \sum_{i=1}^n (h_0(x^{(i)}) - y^{(i)})^2$$

Goal: minimize  $J(\theta_0, \theta_1)$

Linear regression with one variable — Gradient descent

$$J(\theta_0, \theta_1) \longrightarrow \text{want } \min_{\theta_0, \theta_1} J(\theta_0, \theta_1)$$

Outline:

start with some  $\theta_0, \theta_1$

keep changing  $\theta_0, \theta_1$  to reduce  $J(\theta_0, \theta_1)$  until hopefully  
end up with a minimum

repeat until convergence: {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \text{ for } j=0 \text{ and } j=1 \}$$

$\alpha$ : learning rate

Simultaneously update  $\theta_0, \theta_1$

Gradient descent for linear regression

$$\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) = \frac{\partial}{\partial \theta_j} \frac{1}{2n} \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)})^2 = \frac{\partial}{\partial \theta_j} \frac{1}{2n} \sum_{i=1}^n (\theta_0 + \theta_1 x^{(i)} - y^{(i)})^2$$

$$\begin{cases} \frac{\partial}{\partial \theta_0} J(\theta_0, \theta_1) = \frac{1}{n} \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)}) \\ \frac{\partial}{\partial \theta_1} J(\theta_0, \theta_1) = \frac{1}{n} \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x^{(i)} \end{cases}$$

$$\begin{cases} \theta_0 := \theta_0 - \alpha \frac{1}{n} \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)}) \\ \theta_1 := \theta_1 - \alpha \frac{1}{n} \sum_{i=1}^n (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x^{(i)} \end{cases} \quad \text{update } \theta_0, \theta_1 \text{ simultaneously}$$

Cost function for linear regression will always be a convex function (Bowl-shaped), 1 global/local min

"Batch" Gradient Descent

"Batch": Each step of gradient descent uses all the training examples.

Linear Algebra Review

shig, I know everything